

Lasso

Brooke Anderson

March 3, 2016

Reminder– best model to date

Add spline with knots at high temperature– on Kaggle: 0.41048.

```
mod_4 <- glm(count ~ year*hour*workingday*season +  
              weather +  
              ns(temp, knots = c(30, 35)),  
              data = train, family = quasipoisson)
```

Lasso with caret

caret lets you fit a lasso model with `glmnet`, which lets you:

- ▶ Optimize on both complexity parameter (λ) and elastic-net mixing parameter (α)
- ▶ Use RMSLE when tuning
- ▶ Fit different flavors of GLM

See [here](#) for more on `glmnet`.

Lasso with caret

Elastic-net mixing parameter:

$$\alpha = \begin{cases} 0 & \text{ridge penalty} \\ 1 & \text{lasso penalty} \end{cases}$$

Non-zero variance predictors

I was getting some warnings about non-zero variance predictors, so I found and removed them using `nearZeroVar`:

```
my_train <- select(train, -datetime, -registered,  
                    - casual)  
my_train <- model.matrix(count ~ year * season *  
                          workingday * hour +  
                          holiday + temp + atemp +  
                          humidity + windspeed +  
                          month + yday + weather,  
                          data = my_train)  
nzv <- nearZeroVar(my_train)  
my_train <- my_train[, -nzv]
```

Non-zero variance predictors

I need to do the same thing with the testing data (notice I'm using the nzv vector I measured from the training data, not doing a new one for the testing data):

```
my_test <- select(test, -datetime)
my_test <- model.matrix( ~ 1 + year * season *
                        workingday * hour +
                        holiday + temp + atemp +
                        humidity + windspeed +
                        month + yday + weather,
                        data = my_test)
my_test <- my_test[, -nzv]
```


Linear regression with glmnet

```
rmsle_fun <- function(data, lev = NULL,  
                      model = NULL, ...){  
  data$pred[data$pred < 0] <- 0  
  log_p_1 <- log(data$pred + 1)  
  log_a_1 <- log(data$obs + 1)  
  sle <- (log_p_1 - log_a_1)^2  
  rmsle <- sqrt(mean(sle))  
  names(rmsle) <- "rmsle"  
  return(rmsle)  
}  
  
fitControl <- trainControl(method = "cv",  
                           number = 5,  
                           summaryFunction = rmsle_fun)
```

Linear regression with glmnet

```
mod_1 <- train(y = train$count,  
               x = my_train,  
               preProcess = c("center", "scale"),  
               method = "glmnet",  
               trControl = fitControl,  
               metric = "rmsle",  
               maximize = FALSE,  
               family = "gaussian",  
               tuneLength = 5)
```


Linear regression with glmnet



Linear regression with glmnet

```
rmsle <- function(train_preds, actual_preds){  
  train_preds[train_preds < 0] <- 0  
  log_p_1 <- log(train_preds + 1)  
  log_a_1 <- log(actual_preds + 1)  
  sle <- (log_p_1 - log_a_1)^2  
  rmsle <- sqrt(mean(sle))  
  return(rmsle)  
}  
  
train_preds <- predict(mod_1, newdata = my_train)  
train_preds[train_preds < 0] <- 0  
rmsle(train_preds, train$count)
```


```
## [1] 1.388928
```

Linear regression with glmnet

```
test_preds <- predict(mod_1, newdata = my_test)
test_preds[test_preds < 0] <- 0
write_test_preds(test_preds,
                  mod = "elastic_net_gaussian")
```

On Kaggle, this got 1.39085.

Linear regression with glmnet



Poisson GLM with glmnet

From Hastie and Qian:

“The log-likelihood for observations $\{\mathbf{x}_i, y_i\}_1^N$ is given by:

$$l(\beta|X, Y) = \sum_{i=1}^N (y_i(\beta_0 + \beta' \mathbf{x}_i) - e^{\beta_0 + \beta' \mathbf{x}_i})$$

Poisson GLM with glmnet

“As before, we optimize the penalized log-likelihood:

$$\min_{\beta_0, \beta} -\frac{1}{N} l(\beta | X, Y) + \lambda \left((1 - \alpha) \sum_{i=1}^N \beta_i^2 / 2 + \alpha \sum_{i=1}^N |\beta_i| \right)$$

Poisson GLM with glmnet

“Glmnet uses an outer Newton loop, and an inner weighted least-squares loop (as in logistic regression) to optimize this criterion.”

Poisson GLM with glmnet

You need a new rmsle_fun function definition (notice you need to take the exponential of the predicted values):


```
rmsle_fun <- function(data, lev = NULL,
                      model = NULL, ...){
  log_p_1 <- log(exp(data$pred) + 1)
  log_a_1 <- log(data$obs + 1)
  sle <- (log_p_1 - log_a_1)^2
  rmsle <- sqrt(mean(sle))
  names(rmsle) <- "rmsle"
  return(rmsle)
}
```


Poisson GLM with glmnet

Then train the model:

```
mod_2 <- train(y = train$count,  
               x = my_train,  
               preProcess = c("center", "scale"),  
               method = "glmnet",  
               trControl = fitControl,  
               metric = "rmsle",  
               maximize = FALSE,  
               family = "poisson",  
               tuneLength = 5)
```

Poisson GLM with glmnet



Poisson GLM with glmnet

```
rmsle <- function(train_preds, actual_preds){  
  log_p_1 <- log(train_preds + 1)  
  log_a_1 <- log(actual_preds + 1)  
  sle <- (log_p_1 - log_a_1)^2  
  rmsle <- sqrt(mean(sle))  
  return(rmsle)  
}  
  
train_preds <- predict(mod_2, newdata = my_train)  
train_preds <- exp(train_preds)  
rmsle(train_preds, train$count)  
  
## [1] 1.311138
```

Poisson GLM with glmnet

```
test_preds <- predict(mod_2, newdata = my_test)
test_preds <- exp(test_preds)
write_test_preds(test_preds,
                  mod = "elastic_net_poisson")
```

The score on Kaggle was 1.37032.

Poisson GLM with glmnet

